

# Implementation Assignment 2

Rong Yu and Finn Womack


April 21, 2017

## Part 1: Loading Data and Implementing Gradient Batch

After loading the testing and training data into feature and response matrices we implemented the batch gradient decent algorithm with the following learning rates:

$$Rates = \begin{bmatrix} 2 \cdot 10^{-6} \\ 1 \cdot 10^{-6} \\ 3 \cdot 10^{-7} \\ 2 \cdot 10^{-7} \\ 1 \cdot 10^{-7} \end{bmatrix} \quad (1)$$

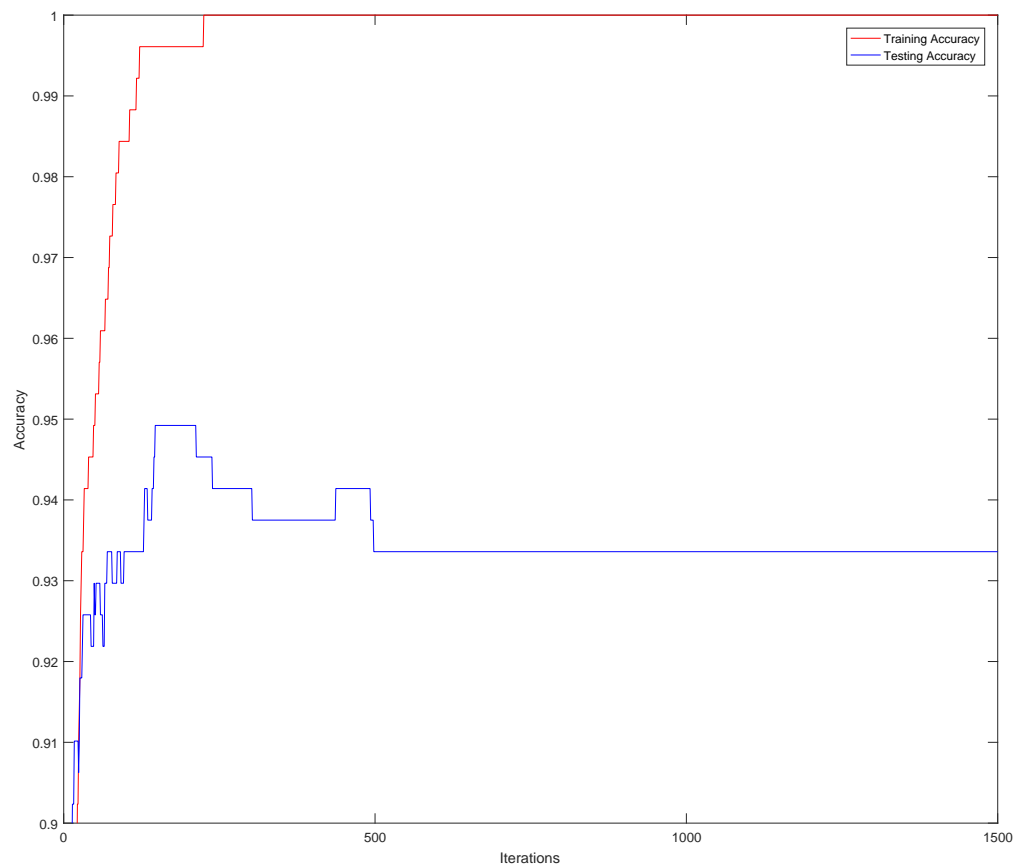
Then after running the algorithm on the above rates we plotted the iterations onto the loss function to get a sense of the convergence rate for the different learning rates:



The learning rate of  $2 \cdot 10^{-6}$  gives the fastest convergence. When we picked larger learning rates we started to get oscillations.

## Part 2: Testing and Training Accuracies

We then selected the learning rate of  $1 \cdot 10^{-7}$  to test the effects of iterations on the testing and training accuracies:



As we can see the training accuracy converges to 1 and the testing accuracy peaks at around 175 iterations and then decreases.

### Part 3: Deriving the Regularization Term

Consider the new objective function:

$$L(w) = \sum_{i=1}^n l(w^T x_i, y_i) + \frac{\lambda}{2} \|W\|_2^2$$

This is the the same as the original objective function with an added term and since the gradient of a sum of functions is the sum of the gradients all we need to do is find the gradient of  $\frac{\lambda}{2} \|W\|_2^2$ :

$$\nabla \frac{\lambda}{2} \|W\|_2^2 = \frac{\lambda}{2} \nabla \|W\|_2^2 \quad (2)$$

$$= \frac{\lambda}{2} \nabla \sum_{i=1}^m w_i^2 \quad (3)$$

$$= \frac{\lambda}{2} \begin{bmatrix} \frac{\partial}{\partial w_1} \sum_{i=1}^m w_i^2 \\ \frac{\partial}{\partial w_2} \sum_{i=1}^m w_i^2 \\ \vdots \\ \frac{\partial}{\partial w_m} \sum_{i=1}^m w_i^2 \end{bmatrix} \quad (4)$$

$$= \frac{\lambda}{2} \begin{bmatrix} 2w_1 \\ 2w_2 \\ \vdots \\ 2w_m \end{bmatrix} \quad (5)$$

$$= \lambda W \quad (6)$$

Thus the new batch gradient code would be as follows:

Given: training examples  $(x_i, y_i), i = 1, \dots, N$

$W \leftarrow [0, 0, \dots, 0]$

Repeat until convergence:

$d \leftarrow [0, 0, \dots, 0]$

For  $i = 1$  to  $N$  do

$\hat{y}_i \leftarrow \frac{1}{1+e^{-w \cdot x_i}}$

$error = y_i - \hat{y}_i$

$d = d + error \cdot (x_i + 2w)$

$w \leftarrow w + \eta d$

## Part 4: Implementing Regularization

To examine the effect of regularization on the algorithm we plotted the testing and training accuracies against the following choices of  $\lambda$ :

$$\lambda \begin{bmatrix} 10^{-6} \\ 10^{-4} \\ 10^{-2} \\ 1 \\ 10^2 \\ 10^4 \\ 10^6 \end{bmatrix} \quad (7)$$

The resulting plot is as follows:

